



Accident severity prediction using the classification algorithm for Seattle city

Coursera Capstone Project

Business understanding

- On daily basis, hundreds of thousands of people travel from A to B using road infrastructure
- One of the less fortunate side or aspect of road transportation is road accidents
- Road accidents not only cause injuries and deaths but also results in loss or damage of properties worth millions of Dollars
- Road accidents are the result of many factors
- In this project, we try to predict the severity of the accident given the road condition, weather condition, type of accident, and many more influencing factors. To effectively predict the severity, we will be using the machine learning algorithms
- The stack holders that can be benefited from an effective ML model which can predict the accident severity are:
 1. People who are commuters and would like to plan their travel
 2. The government regulatory bodies who, can take precautionary steps, to avoid or to reduce the fatality of an accident

Data Understanding

- The data set used for the project is procured from the Seattle Department of Transport and consists of the vehicle accident information
- The data set is in .csv format with 38 columns and a total of 194673 entries or rows
- The label for the selected .csv data set is the column 'SEVERITYCODE'
- Each row in the data set has given a severity code which is either '1' or '2' meaning 'property damage only collision' and 'injury collision', respectively
- By examining the number of 1s and 2s in the label of the given data set, we see that the data is having an imbalance i.e. there are more rows with 1s compared to 2s
- The data must be cleaned in such a way that it avoids the irrelevant attributes and increase the accuracy of the machine learning model without any bias

Data Cleaning


- Before feeding the data to a machine learning model, it should be cleaned or wrangled to avoid problems and false predictions
- The data cleaning or wrangling has been carried out in the following steps:
 1. Removing columns with higher number of missing values
 2. Removing irrelevant columns
 3. Removing columns with repeating and inconsistent data
- After data cleaning we are left with following columns:

| SEVERITYCODE | ADDRTYPE | COLLISIONTYPE | PERSONCOUNT | PEDCOUNT | PEDCYLCOUNT | VEHCOUNT | JUNCTIONTYPE | WEATHER | ROADCOND | LIGHTCOND | HITPARKEDCAR |
|--------------|--------------|---------------|-------------|----------|-------------|----------|---|----------|----------|-------------------------|--------------|
| 2 | Intersection | Angles | 2 | 0 | 0 | 2 | At Intersection (intersection related) | Overcast | Wet | Daylight | N |
| 1 | Block | Sideswipe | 2 | 0 | 0 | 2 | Mid-Block (not related to intersection) | Raining | Wet | Dark - Street Lights On | N |
| 1 | Block | Parked Car | 4 | 0 | 0 | 3 | Mid-Block (not related to intersection) | Overcast | Dry | Daylight | N |
| 1 | Block | Other | 3 | 0 | 0 | 3 | Mid-Block (not related to intersection) | Clear | Dry | Daylight | N |
| 2 | Intersection | Angles | 2 | 0 | 0 | 2 | At Intersection (intersection related) | Raining | Wet | Daylight | N |

Data Formatting

- After the data is cleaned, the categorical values and numerical values are separated into two different variables to convert only the categorical attributes. After that, the unbalance in the data set must be corrected
- One Hot Encoding: it is a method in which the categorical values are automatically converted into numeric values

```
df_dummies = pd.get_dummies(df[["ADDRTYPE", "COLLISIONTYPE", "JUNCTIONTYPE", "WEATHER", "ROADCOND", "LIGHTCOND", "HITPARKEDCAR"]])
```



```
df['SEVERITYCODE'].value_counts()  
1    95921  
2    49448  
Name: SEVERITYCODE, dtype: int64
```

- As we discussed, our data set has an unbalance in the labelled attribute 'SEVERITYCODE', which if present, can result in a bias while predicting the unknown values. The balance in the data is achieved by making the number 1s the same as 2s
- The balance is achieved by using the NearMiss package of imblearn library

Number of columns after data cleaning

- The figure on the right shows the columns of independent variable after one hot encoding


```
Pre_X.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 145369 entries, 0 to 194672  
Data columns (total 47 columns):
```

| # | Column | Non-Null Count | Dtype |
|----|--|-----------------|-------|
| 0 | PERSONCOUNT | 145369 non-null | int64 |
| 1 | PEDCOUNT | 145369 non-null | int64 |
| 2 | PEDCYLCOUNT | 145369 non-null | int64 |
| 3 | VEHCOUNT | 145369 non-null | int64 |
| 4 | ADDRTYPE_Alley | 145369 non-null | uint8 |
| 5 | ADDRTYPE_Block | 145369 non-null | uint8 |
| 6 | ADDRTYPE_Intersection | 145369 non-null | uint8 |
| 7 | COLLISIONTYPE_Angles | 145369 non-null | uint8 |
| 8 | COLLISIONTYPE_Cycles | 145369 non-null | uint8 |
| 9 | COLLISIONTYPE_Head On | 145369 non-null | uint8 |
| 10 | COLLISIONTYPE_Left Turn | 145369 non-null | uint8 |
| 11 | COLLISIONTYPE_Parked Car | 145369 non-null | uint8 |
| 12 | COLLISIONTYPE_Pedestrian | 145369 non-null | uint8 |
| 13 | COLLISIONTYPE_Rear Ended | 145369 non-null | uint8 |
| 14 | COLLISIONTYPE_Right Turn | 145369 non-null | uint8 |
| 15 | COLLISIONTYPE_Sideswipe | 145369 non-null | uint8 |
| 16 | JUNCTIONTYPE_At Intersection (but not related to intersection) | 145369 non-null | uint8 |
| 17 | JUNCTIONTYPE_At Intersection (intersection related) | 145369 non-null | uint8 |
| 18 | JUNCTIONTYPE_Driveway Junction | 145369 non-null | uint8 |
| 19 | JUNCTIONTYPE_Mid-Block (but intersection related) | 145369 non-null | uint8 |
| 20 | JUNCTIONTYPE_Mid-Block (not related to intersection) | 145369 non-null | uint8 |
| 21 | JUNCTIONTYPE_Ramp Junction | 145369 non-null | uint8 |
| 22 | WEATHER_Blowing Sand/Dirt | 145369 non-null | uint8 |
| 23 | WEATHER_Clear | 145369 non-null | uint8 |
| 24 | WEATHER_Fog/Smog/Smoke | 145369 non-null | uint8 |
| 25 | WEATHER_Overcast | 145369 non-null | uint8 |
| 26 | WEATHER_Partly Cloudy | 145369 non-null | uint8 |
| 27 | WEATHER_Raining | 145369 non-null | uint8 |
| 28 | WEATHER_Severe Crosswind | 145369 non-null | uint8 |
| 29 | WEATHER_Sleet/Hail/Freezing Rain | 145369 non-null | uint8 |
| 30 | WEATHER_Snowing | 145369 non-null | uint8 |
| 31 | ROADCOND_Dry | 145369 non-null | uint8 |
| 32 | ROADCOND_Ice | 145369 non-null | uint8 |
| 33 | ROADCOND_Oil | 145369 non-null | uint8 |
| 34 | ROADCOND_Sand/Mud/Dirt | 145369 non-null | uint8 |
| 35 | ROADCOND_Snow/Slush | 145369 non-null | uint8 |
| 36 | ROADCOND_Standing Water | 145369 non-null | uint8 |
| 37 | ROADCOND_Wet | 145369 non-null | uint8 |
| 38 | LIGHTCOND_Dark - No Street Lights | 145369 non-null | uint8 |
| 39 | LIGHTCOND_Dark - Street Lights Off | 145369 non-null | uint8 |
| 40 | LIGHTCOND_Dark - Street Lights On | 145369 non-null | uint8 |
| 41 | LIGHTCOND_Dark - Unknown Lighting | 145369 non-null | uint8 |
| 42 | LIGHTCOND_Dawn | 145369 non-null | uint8 |
| 43 | LIGHTCOND_Daylight | 145369 non-null | uint8 |
| 44 | LIGHTCOND_Dusk | 145369 non-null | uint8 |
| 45 | HITPARKEDCAR_N | 145369 non-null | uint8 |
| 46 | HITPARKEDCAR_Y | 145369 non-null | uint8 |

```
dtypes: int64(4), uint8(43)  
memory usage: 11.5 MB
```

Exploratory Data Analysis



- Before creating and splitting the data, we can explore our clean and balanced data for the correlation between different variables
- The values shown are between 1 and -1. If the value is near zero, it means that the correlation is weak. Values close to 1 and -1 shows a strong positive or negative correlation, respectively
- From the correlation matrix, we can see a strong positive correlation between the 'ROADCOND_Wet' and 'WEATHERCOND_Rainy', 'ADDRTYPE_Intersection', and 'JUNCTIONTYPE_At Intersection', and 'PEDCOUNT' and 'COLLISIONTYPE_Pedestrian'
- Moreover, a strong negative correlation is observed between the attributes like 'WEATHER_Raining' and 'ROADCOND_Dry' as the given condition is contradictory

Data Normalization

- The data should be standardized to reduce the error in the prediction because of large values in the columns. The standardization only applies to the column with the numerical value, not to the already encoded categorical values
- But before carrying out data standardization, the data set must be split into testing and training data set

```
X_train, X_test, y_train, y_test = train_test_split( X, y, test_size=0.2, random_state=4)
print ('Train set:', X_train.shape, y_train.shape)
print ('Test set:', X_test.shape, y_test.shape)
```

```
Train set: (79116, 47) (79116,)
Test set: (19780, 47) (19780,)
```

- We can use StandardScaler() from the preprocessing module of the Scikit Learn library. First, we create a StandardScaler() object and then apply .fit_transform method on the numerical columns of our testing and training data set separately

```
scaler = StandardScaler()
X_train[['PERSONCOUNT', 'PEDCOUNT', 'PEDCYLCOUNT', 'VEHCOUNT']] = scaler.fit_transform(X_train[['PERSONCOUNT', 'PEDCOUNT', 'PEDCYLCOUNT', 'VEHCOUNT']])

scaler = StandardScaler()
X_test[['PERSONCOUNT', 'PEDCOUNT', 'PEDCYLCOUNT', 'VEHCOUNT']] = scaler.fit_transform(X_test[['PERSONCOUNT', 'PEDCOUNT', 'PEDCYLCOUNT', 'VEHCOUNT']])
```


Modeling

- Our goal is to predict the severity of an accident in binary classes i.e. the severity could be either 1 or 2
- This clearly shows that this problem can be solved by using Classification machine learning algorithms as the target attribute is a categorical variable
- Classification is a supervised learning approach in which the label of our data is categorized into a discrete set of categories or classes. Classification determines a class label for an unlabelled case
- The following classification models are used :
 1. KNN model
 2. Decision Tree model
 3. SVM Model
 4. Logistic regression model

KNN Model


- k-nearest neighbour is an algorithm that takes several labelled points and uses them to label other points based on the similarity to other cases
- Initially, a value of $k = 4$ is assumed and the distance of unknown cases is calculated. Then k-observations in the training data that are nearest to the unknown data point is selected. Afterward, the response of the unknown data point using the most popular response value from the k-nearest neighbours is predicted

```
k = 4  
neigh = KNeighborsClassifier(n_neighbors = k).fit(X_train,y_train)  
neigh
```

```
KNeighborsClassifier(n_neighbors=7)
```

```
yhat_knn = neigh.predict(X_test)  
yhat_knn[0:5]
```

```
array([2, 1, 2, 1, 2], dtype=int64)
```



- Now, to calculate the optimum value of k , we can loop different values of k using for loop and then visualize the value of k vs. the accuracy score. The visualization is shown below. After the evaluation, we select $k = 7$ for our model

Decision Tree Model

- The basic intuition behind a decision tree is to map out all the possible decision paths in the form of a tree
- First, the attributes are considered and chosen one after another, then the significance of the attributes in the splitting of data is calculated
- Thereafter, the data based on the value of the best attribute is split. Finally, this process is repeated until all the attributes are split into the leaf nodes
- The aim here is to increase the accuracy as much as possible and keep the entropy smallest

```
DecTree = DecisionTreeClassifier(criterion="entropy", max_depth = 4)
```

```
DecTree.fit(X_train, y_train)
```

```
DecisionTreeClassifier(criterion='entropy', max_depth=4)
```

```
yhat_DecTree = DecTree.predict(X_test)  
yhat_DecTree[0:10]
```

```
array([2, 2, 2, 2, 2, 2, 1, 2, 2, 2], dtype=int64)
```

SVM Model

- The support vector machine algorithm classifies cases by finding a separator
- SVM works by mapping data to a high-dimensional feature space in such a way that data points can be classified even when the data is not linearly separable in 2D space
- When plotting the data in multi-dimensional space, the category or class or the data points are separated using a separator represented as a hyperplane. The mapping of the data in higher dimensions is called kernelling. When creating an SVM object clf, RBF (Random Basis Function) type of kernel is been selected

```
clf = svm.SVC(kernel='rbf')
```

```
clf.fit(X_train, y_train)
```

```
SVC()
```

```
yhat_clf = clf.predict(X_test)  
yhat_clf [0:5]
```

```
array([2, 2, 1, 2, 1], dtype=int64)
```

Logistic Regression Model

- Logistic regression can be used to determine the probability of the outcome and to understand the impact of the feature
- The probability is predicted by fitting the data using a sigmoid or logistic function. After fitting the data, the model output *yhat* and the actual output of severity are compared, and the error value is recorded
- The aim here is to reduce the cost i.e. maximizing the accuracy by constantly changing the value of theta. In our instance of logistic regression, the value of C (indicates inverse of regularization strength) is taken as 0.01 (must be a positive float). The numerical optimizer used is 'liblinear'

```
LR = LogisticRegression(C=0.01, solver='liblinear').fit(X_train,y_train)
```

```
yhat_LR = LR.predict(X_test)  
yhat_LR[0:5]
```

```
array([2, 2, 1, 2, 1], dtype=int64)
```

Model Evaluation

- For the evaluation of the models, we have selected 3 types of evaluation criteria
- All the implemented method gives a score between 0 and 1 where 0 being least accurate and 1 being very accurate
- The following methods are used :
 1. Accuracy Scores
 2. Jaccard Similarity Scores
 3. F1 Score

Result and Conclusion

- The following table shows the above-mentioned accuracy scores for 4 of our selected models

| Algorithm | Jaccard | F1-score | Train set Accuracy | Test set Accuracy |
|----------------------------|---------|----------|--------------------|-------------------|
| KNN | 0.4884 | 0.6832 | 0.7004 | 0.6853 |
| Decision Tree | 0.4173 | 0.6679 | 0.6922 | 0.6861 |
| SVM | 0.4943 | 0.6981 | 0.7125 | 0.7023 |
| Logistic Regression | 0.4736 | 0.6815 | 0.6960 | 0.6860 |

- Here, we can see that no algorithm has resulted in an accuracy of more than 71%. Considering the result, we can conclude that most model has shown the score in very close range of approx. $\pm 2\%$. The Support Vector Machine provides the highest accuracy scores in all the categories. Thus, for this project SVM classifiers proves to be the relatively most accurate model to predict the severity of the accident given different attributes or conditions
- Moreover, it has also been observed that the time required to fit the model and predicting the output, Logistic regression, and Decision tree took the least time
- The maximum test accuracy of the SNM model is 70% which can be a result of underfitting and can be improved using more data. Due to an imbalance in the selected data set, many observations are dropped, so in the future, a more balanced data set can be used

Future scopes

- Many important columns such as 'UNDERINFL' and 'INATTENTIONIND' which shows whether the driver is under influence of drugs and driver inattention respectively, has been dropped due to data inconsistency and missing values
- These could have been good factors in predicting accident severity. In the future, we can select more consistent data
- Moreover, we only had a binary label in our data set (1s and 2s), in the future, we can have data with multiclass labels to make machine learning model more challenging